

What is the default number of partition after wide transformation?

Wide transformations in distributed computing involve operations that require shuffling data across partitions. This shuffling is necessary because these transformations need to bring together data from multiple partitions to compute the final result. Examples include `groupByKey()`, `reduceByKey()`, `join()`, `distinct()`, and `intersect()`. They resemble the shuffle-and-sort phase of MapReduce, as they involve rearranging and redistributing data across the cluster before computations. However, they can be costly due to data movement and network overhead, so they should be used judiciously and optimized for performance.

Overall, the default value of 200 partitions strikes a balance between parallelism, memory overhead, task granularity, and shuffle efficiency. However, it's important to note that the optimal number of partitions can vary depending on factors such as the size of the dataset, the available resources (e.g., CPU cores, memory), and the nature of the computations being performed. In practice, it's often necessary to adjust the number of partitions based on the specific requirements and characteristics of each Spark application.

In the program I have converted the dataframe into RDD to access the number of partition. When you work with RDDs directly, you have access to the number of partitions, and you can perform operations based on this information.

However, when you work with DataFrames, which are a higher-level abstraction built on top of RDDs, the concept of partitions is abstracted away. DataFrames are designed to provide a more structured and efficient way of working with distributed data, and they encapsulate many optimization techniques, including partitioning, under the hood.

CODE:

```
import org.apache.spark.sql.{SparkSession, DataFrame}

object DefaultPartition {

  def main (args: Array [String]): Unit = {

    // Create SparkSession
    val spark = SparkSession.builder()
      .appName("dataframepartition")
      .master("local[*]") // Run locally using all available cores
      .getOrCreate()
```

```
// Sample data for demonstration
val data = Seq(
  ("A", 34),
  ("B", 45),
  ("C", 28),
  ("D", 56),
  ("E", 40)
)

// Create a DataFrame from the sample data
import spark.implicits._
val df = data.toDF("Name", "Age")

// Check the number of partitions before shuffle
val initialPartitions = df.rdd.getNumPartitions
println(s"Number of partitions before shuffle: $initialPartitions")

// Perform a shuffle operation (group by)
val groupedDF: DataFrame = df.groupBy("Age").count()

// Check the number of partitions after shuffle
val partitionsAfterShuffle = groupedDF.rdd.getNumPartitions
println(s"Number of partitions after shuffle: $partitionsAfterShuffle")

// Show the result
groupedDF.show()

val delayInMilliseconds = 5 * 60 * 1000 // 5 minutes in milliseconds
println(s"Pausing execution for ${delayInMilliseconds / 1000 / 60}
minutes before shutting down Spark session...")
Thread.sleep(delayInMilliseconds)

// Stop SparkSession
spark.stop()
}
```

OUTPUT :

24/02/16 09:38:46 INFO CodeGenerator: Code generated in 22.0292 ms

Number of partitions before shuffle: 4

24/02/16 09:38:47 INFO CodeGenerator: Code generated in 48.4815 ms

24/02/16 09:38:47 INFO CodeGenerator: Code generated in 67.2185 ms

24/02/16 09:38:47 INFO CodeGenerator: Code generated in 13.7245 ms

Number of partitions after shuffle: 200

24/02/16 09:38:48 INFO CodeGenerator: Code generated in 37.4185 ms

EXPLANATION OF THE OUTPUT:

It appears that the number of partitions before the shuffle operation is 4, and after the shuffle operation, it becomes 200. This behavior is consistent with the default behavior of Spark:

1. **Before Shuffle:** The DataFrame **df** initially has 4 partitions. This number is determined by Spark's default parallelism, which is typically set to the number of available CPU cores on your local machine when running in local mode (**local[*]**).
2. **After Shuffle:** The **groupBy** operation triggers a shuffle, where data is redistributed across partitions based on the grouping key (**Age** in this case). During the shuffle, Spark typically uses the value of **spark.sql.shuffle.partitions**, which defaults to 200, to determine the number of partitions for the resulting DataFrame **groupedDF**. Hence, after the shuffle, **groupedDF** has 200 partitions.

This behavior is expected and is in line with Spark's design to optimize parallelism and performance during shuffle operations. If you want to control the number of partitions explicitly after the shuffle operation, you can use methods like **repartition** or **coalesce** to adjust the number of partitions according to your requirements.

What is persistence partitioning?

Persistence partitioning refers to a technique used in computer science and database management to optimize the storage and retrieval of data. It involves dividing the data into multiple partitions or segments based on certain criteria such as time, range, or specific attributes.

The primary goal of persistence partitioning is to improve performance, scalability, and manageability of the data storage system. By partitioning data, operations like querying, updating, and deleting can be performed more efficiently because they can be targeted to specific partitions rather than having to scan through the entire dataset.

Here are some common types of persistence partitioning:

1. **Time-based partitioning:** Data is partitioned based on time intervals, such as days, months, or years. This is particularly useful for applications dealing with time-series data like logs, sensor data, or financial transactions.
2. **Range partitioning:** Data is partitioned based on a specific range of values, such as numerical ranges or alphabetical ranges. For example, in a database table containing customer information, data might be partitioned based on the first letter of the customer's last name.
3. **Hash partitioning:** Data is partitioned based on a hash function applied to a particular attribute or combination of attributes. This ensures a more even distribution of data across partitions, which can help balance the load on the storage system.
4. **List partitioning:** Data is partitioned based on predefined lists of values for a specific attribute. For example, a database table containing sales data might be partitioned based on the region in which the sale occurred.
5. **Composite partitioning:** Data is partitioned using a combination of multiple partitioning strategies. For instance, a system might use a combination of range and hash partitioning to achieve both data locality and even distribution.